


Question 10.1

In R, you can use the tree package or the rpart package, and the randomForest package. For each model, describe one or two qualitative takeaways you get from analyzing the results (i.e., don't just stop when you have a good model, but interpret it too).


Using the same crime data set uscrime.txt as in Questions 8.2 and 9.1, find the best model you can using

(a) a regression tree model, and

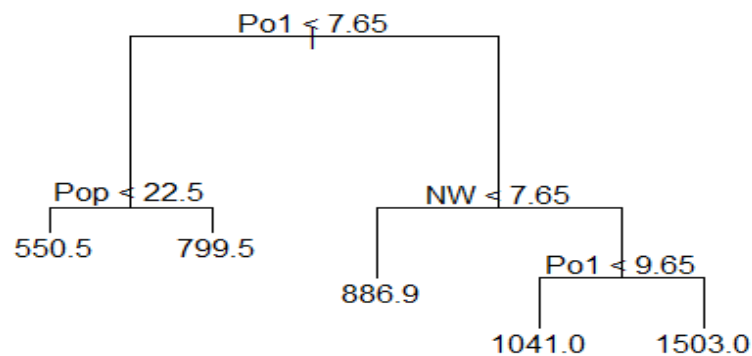
The code for my approach can be found in [appendix 10.1A](#). In my approach, I discovered that there were four significant factors used in branching. They were Po1, Pop, LF, and NW. The regression tree generated is as follows:



As can be seen here, the factors Pop and Po1 are used twice in the tree. I also wanted to investigate further and see if pruning the tree could make a better model. For this, I plotted the tree deviance against the tree size (in terms of terminal nodes):



As can be seen here, it seems that having 5 or 7 terminal nodes produces the least deviance. Because I had already generated the original tree with seven terminal nodes, I decided to create another tree pruned to five nodes:



Here, I see that Po1 appears twice like in the model with seven terminal nodes, but LF is no longer present. Pop no longer appears twice either. To compare both models, I decided to look at the residual mean deviances. Surprisingly, I found that pruning the tree increased the residual mean deviance to 54210 as compared to the 47390 for the original tree with seven terminal nodes. Therefore, it seems that pruning the tree is a wrong decision here. I felt that I should investigate this more too though, and I found the R-squared values for the original tree (0.7245) and the pruned tree (0.6691), which indicates that, again, the original tree is likely better. Finally,

I wanted to use cross-validation to check model quality. As such, I looked at the SSE for each size of the trees without and with cross-validation:

```
#Without CV

prune.tree(data_tree)$size

## [1] 7 6 5 4 3 2 1

prune.tree(data_tree)$dev

## [1] 1895722 2013257 2276670 2632631 3364043 4383406 6880928
```

```
#With CV

cv_data_tree$size

## [1] 7 6 5 4 3 2 1

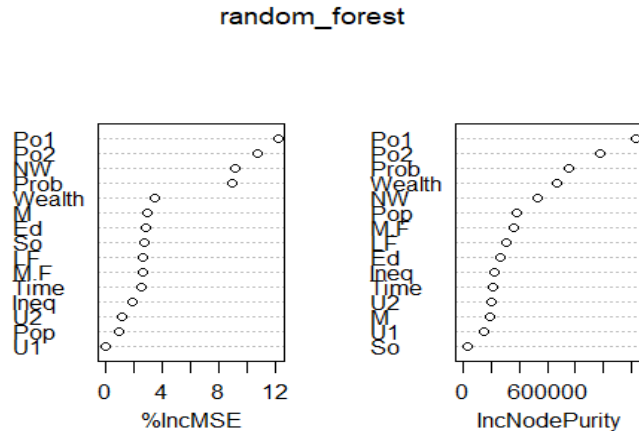
cv_data_tree$dev

## [1] 7049681 7083012 7008434 7402184 7739300 7607425 8773298
```

When I look at this information, the errors are much higher in the cross-validation model, which indicates to me that there is some pretty serious overfitting going on in the original model. Additionally, when I consider the factors in the pruned and unpruned trees, it appears to me that Po1 is an extremely important factor since it is present in both trees and is used several times in branching decisions. However, there is a possibility that LF is less important than the others factors (as it was dropped in the pruned tree).

(b) a random forest model.

The code for this approach can be found in [appendix 10.1B](#). Because of the way random forest works (i.e., generating many trees and then aggregating them), it is able to reduce overfitting. In the regression tree model, I observed that there was a decent bit of overfitting going on, so I investigated the random forest model to see if it would address that. In this process, I obtained an R-squared value of 0.4108 with the random forest. This would make sense if there was less overfitting in the random forest model as compared to the regression tree model. I also investigated the importance of different predictors in the random forest model:



Here, when I compared it to the tree model, I also saw that the random forest model thought the Po1 factor was very important too. Ultimately, I think the random forest model is a good model that prioritizes similar factors while decreasing the risks of overfitting present in the tree model.

Question 10.2

Describe a situation or problem from your job, everyday life, current events, etc., for which a logistic regression model would be appropriate. List some (up to 5) predictors that you might use.

I think a situation where a logistic regression model would be appropriate is if you were investigating how certain variables affect admission into university. For this, the response variable would be a binary one: admittance or rejection. I think some predictors that could be used are: prestige of the student's high school, student's SAT scores, student's ACT scores, length of student's extracurricular involvement, and ethnicity. Using these predictors, a model could be created to predict whether or not a student is likely to be accepted to a university.


Question 10.3

1. Using the GermanCredit data set `germancredit.txt` from <http://archive.ics.uci.edu/ml/machine-learning-databases/statlog/german/> (description at <http://archive.ics.uci.edu/ml/datasets/Statlog+%28German+Credit+Data%29>), use logistic regression to find a good predictive model for whether credit applicants are good credit risks or not. Show your model (factors used and their coefficients), the software output, and the quality of fit. You can use the `glm` function in R. To get a logistic regression (logit) model on data where the response is either zero or one, use `family=binomial(link="logit")` in your `glm` function call.

The code for this question can be found in [appendix 10.3.1](#). Because there were several variables in the data set and not all of them were significant, I wanted to develop the best model possible. In order to do this, I iterated through several regression models choosing only the significant variables from each one. In the end, I was able to develop the following model:

Variable, Value	Coefficient
Intercept	-4.529e-01
V1, A13	-9.802e-01
V1, A14	-1.637e+00
V2	3.102e-02
V3, A34	-6.940e-01
V4, A41	-1.715e+00
V4, A410	-2.311e+00
V4, A42	-6.579e-01
V4, A43	-9.249e-01
V4, A49	-8.619e-01
V5	1.295e-04
V6, A64	-1.294e+00
V6, A65	-7.435e-01
V8	3.301e-01
V10, A103	-1.325e+00
V14, A143	-8.612e-01
V15, A152	-4.660e-01

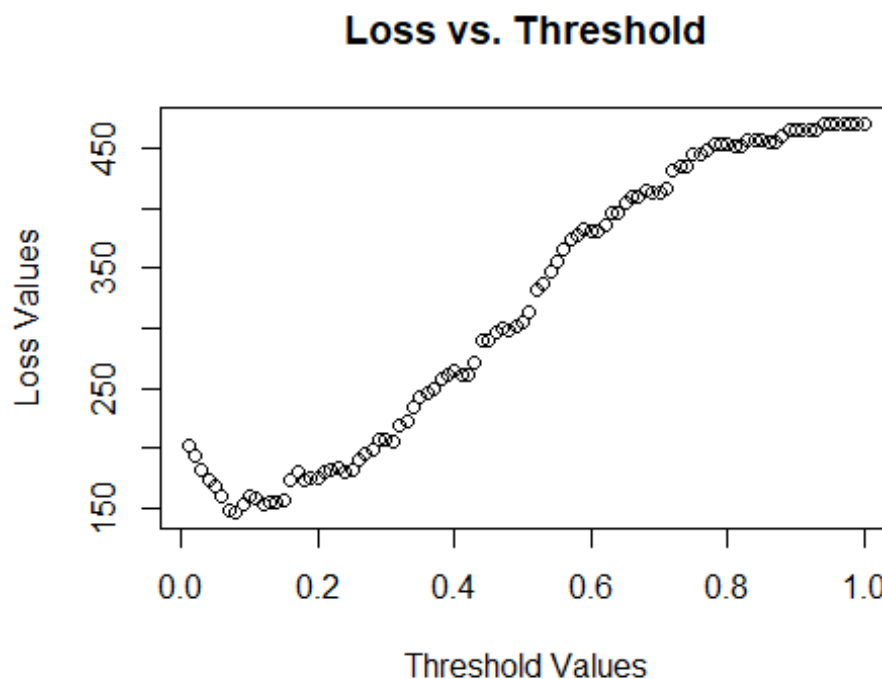
The Akaike Information Criterion value I obtained for this was 683.02. When I validated the model, I obtained an accuracy of 0.7433333 (threshold = 0.5) and a receiver operating characteristic (ROC) curve as follows:



My value for area under the curve (AUC) was 0.6483.

2. Because the model gives a result between 0 and 1, it requires setting a threshold probability to separate between “good” and “bad” answers. In this data set, they estimate that incorrectly identifying a bad customer as good, is 5 times worse than incorrectly classifying a good customer as bad. Determine a good threshold probability based on your model.

The code for this approach can be found in [appendix 10.3.2](#). To determine a good threshold probability, I iterated through different threshold values from (0.01 to 1) and performed cost of loss calculations for each threshold. The plot of the cost of loss against threshold values is as follows:



The specific cost of loss values are as follows (each index is a threshold value; e.g., the first value shows the loss for threshold 0.01; the second shows loss for threshold 0.02, ..., all the way to a threshold of 1):

```
##[1] 202 194 182 173 169 160 149 147 154 160 158 154 156 156 157 173 180 174

##[19] 176 176 180 182 184 181 183 190 195 199 207 207 206 219 222 234 243 246

##[37] 250 258 262 265 262 262 272 290 289 296 300 298 302 305 314 332 337 347

##[55] 356 365 374 377 382 381 381 386 396 395 405 409 409 414 412 412 416 431
```

```
##[73] 435 434 444 444 448 453 453 453 452 451 456 456 456 455 455 460 465 465
```

```
##[91] 465 465 465 470 470 470 470 470 470 470
```

As can be seen, it appears that the lowest loss occurs when threshold is 0.08 (loss = 147). With a threshold of 0.08, the accuracy is 0.5366667 and the AUC of its ROC curve is 0.6568. In general, the losses are relatively low from a threshold of 0.07 (loss = 149) until a threshold of 0.15 (loss = 157) but, after that, they begin to rise relatively steeply. It can also be seen that if we chose a threshold of 0.5 like in question 10.3.1, the loss would be much higher at 305. Therefore, it's important not to just choose the threshold with the highest accuracy but to also consider the cost of loss associated with it.

Appendix

Question 10.1A

#This is 10.1A

```
library(DAAG); library(tree)
```

```
set.seed(7)
```

#Load data

```
data <- read.table("C:\\Users\\User\\OneDrive\\Desktop\\Data 10.1\\uscrime.txt",  
stringsAsFactors = F, header = T)
```

#Create regression tree

```
data_tree <- tree(Crime~., data = data)
```

#Only four factors were used to create the tree (Po1, Pop, LF, and NW)

```
summary(data_tree)
```

```
##
```

```
## Regression tree:
```

```
## tree(formula = Crime ~ ., data = data)
```

```
## Variables actually used in tree construction:
```

```
## [1] "Po1" "Pop" "LF" "NW"
```

```
## Number of terminal nodes: 7
```

```
## Residual mean deviance: 47390 = 1896000 / 40
```

```
## Distribution of residuals:
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
## -573.900 -98.300  -1.545    0.000 110.600  490.100
```

```
data_tree$frame
```

```
##      var  n      dev      yval splits.cutleft splits.cutright  
## 1    Po1 47 6880927.66  905.0851      <7.65      >7.65  
## 2    Pop 23 779243.48  669.6087      <22.5      >22.5  
## 4     LF 12 243811.00  550.5000      <0.5675     >0.5675  
## 8 <leaf> 7  48518.86  466.8571  
## 9 <leaf> 5  77757.20  667.6000  
## 5 <leaf> 11 179470.73  799.5455  
## 3     NW 24 3604162.50 1130.7500      <7.65      >7.65  
## 6    Pop 10 557574.90  886.9000      <21.5      >21.5  
## 12 <leaf> 5 146390.80 1049.2000  
## 13 <leaf> 5 147771.20  724.6000  
## 7    Po1 14 2027224.93 1304.9286      <9.65      >9.65  
## 14 <leaf> 6 170828.00 1041.0000  
## 15 <leaf> 8 1124984.88 1502.8750
```



```

#Plot the tree
plot(data_tree)
text(data_tree)

#We can consider the deviance of trees with different numbers of terminal nodes
#Based on the values, we can decide how to prune the tree if we want

cv_data_tree <- cv.tree(data_tree)
plot(cv_data_tree$size, cv_data_tree$dev, type= "b")

#Based on the plot, the lowest deviation is with 5 or 7 terminal nodes

#Pruning the tree
terminal_nodes <- 5
prune_data_tree <- prune.tree(data_tree, best = terminal_nodes)
plot(prune_data_tree)
text(prune_data_tree)

summary(prune_data_tree)

##
## Regression tree:
## snip.tree(tree = data_tree, nodes = c(4L, 6L))
## Variables actually used in tree construction:
## [1] "Po1" "Pop" "NW"
## Number of terminal nodes: 5
## Residual mean deviance: 54210 = 2277000 / 42
## Distribution of residuals:
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -573.9 -107.5   15.5     0.0  122.8   490.1

#If we compare the residual mean deviance, pruning the tree increased it
# from 47390 to 54210
#Therefore, let's stick with the unaltered model

#Let's calculate the quality of fit of the model
data_tree_pred <- predict(data_tree, data=data[,1:15])
RSS <- sum((data_tree_pred - data[,16])^2)
TSS <- sum((data[,16] - mean(data[,16]))^2)
R2 <- 1 - RSS/TSS
R2

## [1] 0.7244962

#The R-squared is therefore 0.7245

#We can also investigate the R-squared value if we used the pruned tree
data_tree_pred2 <- predict(prune_data_tree, data=data[,1:15])
RSS <- sum((data_tree_pred2 - data[,16])^2)
TSS <- sum((data[,16] - mean(data[,16]))^2)

```

```

R2 <- 1 - RSS/TSS
R2

## [1] 0.6691333

#We see that it is lower than the R-squared of the unaltered model

#As we used the training data above, we can also use CV to check model quality
#We can check the SSE for each size of tree without cross-validation

prune.tree(data_tree)$size

## [1] 7 6 5 4 3 2 1

prune.tree(data_tree)$dev

## [1] 1895722 2013257 2276670 2632631 3364043 4383406 6880928

#Let's check the cross validation results now
cv_data_tree$size

## [1] 7 6 5 4 3 2 1

cv_data_tree$dev

## [1] 7049681 7083012 7008434 7402184 7739300 7607425 8773298

#These errors are much, much larger, which indicates overfitting in original model

```

Question 10.1B

```

#This is 10.1B

library(DAAG); library(randomForest)

## randomForest 4.7-1.1

## Type rfNews() to see new features/changes/bug fixes.

set.seed(8)

#Load data
data <- read.table("C:\\Users\\User\\OneDrive\\Desktop\\Data 10.1\\uscrime.txt",
  stringsAsFactors = F, header = T)

#Generate the random forest
ntry <- 4
random_forest <- randomForest(Crime~., data = data, mtry = ntry, importance=T
RUE)
summary(random_forest)

```

```
##           Length Class  Mode
## call           5   -none- call
## type           1   -none- character
## predicted      47   -none- numeric
## mse            500   -none- numeric
## rsq            500   -none- numeric
## oob.times      47   -none- numeric
## importance      30   -none- numeric
## importanceSD    15   -none- numeric
## localImportance 0    -none- NULL
## proximity       0    -none- NULL
## ntree           1    -none- numeric
## mtry            1    -none- numeric
## forest         11   -none- list
## coefs           0    -none- NULL
## y              47   -none- numeric
## test           0    -none- NULL
## inbag           0    -none- NULL
## terms           3    terms  call

random_forest

##
## Call:
## randomForest(formula = Crime ~ ., data = data, mtry = ntry, importance =
TRUE)
##           Type of random forest: regression
##           Number of trees: 500
## No. of variables tried at each split: 4
##
##           Mean of squared residuals: 86264.36
##           % Var explained: 41.08

#Check quality of model
pred_data <- predict(random_forest)
RSS <- sum((pred_data - data[,16])^2)
TSS <- sum((data[,16] - mean(data[,16]))^2)
R2 <- 1 - RSS/TSS
R2

## [1] 0.4107735

#Look at importance of the model
importance(random_forest)

##           %IncMSE IncNodePurity
## M           2.96976642      187734.53
## So           2.74235655       27134.61
## Ed           2.87314506       263528.42
## Po1          12.15681006      1229751.02
## Po2          10.72902168       974488.89
```

```
## LF      2.67065010      306597.09
## M.F     2.65719740      358124.59
## Pop     1.00552773      379656.81
## NW      9.17355993      526536.45
## U1      0.04785962      143931.84
## U2      1.15847648      205050.56
## Wealth  3.44995938      672172.18
## Ineq    1.94745173      226982.78
## Prob    8.93316256      758513.30
## Time    2.58330174      213746.67
```

```
varImpPlot(random_forest)
```

Question 10.3.1

#Load Libraries

```
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      cov, smooth, var
```

#Load Data

```
data <- read.table("C:\\Users\\User\\OneDrive\\Desktop\\Data 10.3\\germancred
it.txt",
```

```
stringsAsFactors = F, header = F)
```

#Convert the 1s and 2s to 0s and 1s for the logistic regression

```
data$V21[data$V21==1] <- 0
```

```
data$V21[data$V21==2] <- 1
```

```
set.seed(10)
```

#Use a 70-30 split of training and testing data

```
nrows <- nrow(data)
```

```
train_set <- sample(1:nrows, size = round(nrows*0.7))
```

```
train <- data[train_set,]
```

```
validate <- data[-train_set,]
```

#Perform iterations to create logistic regression model

#Use all variables first

```
lreg <- glm(V21~., family=binomial(link = "logit"), data = train)
```

```
summary(lreg)
```

```
##
```

```
## Call:
```

```
## glm(formula = V21 ~ ., family = binomial(link = "logit"), data = train)
```

```

##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2454  -0.6866  -0.3350   0.6271   2.7315
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  3.709e-01  1.388e+00   0.267 0.789277
## V1A12        -3.518e-02  2.664e-01  -0.132 0.894918
## V1A13        -1.025e+00  4.638e-01  -2.209 0.027150 *
## V1A14        -1.598e+00  2.848e-01  -5.611 2.01e-08 ***
## V2           3.026e-02  1.143e-02   2.648 0.008091 **
## V3A31        -4.026e-01  6.930e-01  -0.581 0.561304
## V3A32        -6.337e-01  5.518e-01  -1.148 0.250800
## V3A33        -1.038e+00  5.980e-01  -1.735 0.082652 .
## V3A34        -1.439e+00  5.609e-01  -2.565 0.010308 *
## V4A41        -1.865e+00  4.660e-01  -4.003 6.26e-05 ***
## V4A410       -2.376e+00  1.067e+00  -2.227 0.025961 *
## V4A42        -1.008e+00  3.243e-01  -3.110 0.001874 **
## V4A43        -1.041e+00  3.035e-01  -3.429 0.000606 ***
## V4A44        -8.000e-01  1.360e+00  -0.588 0.556372
## V4A45        -1.948e-01  6.255e-01  -0.311 0.755457
## V4A46        -9.934e-02  4.589e-01  -0.216 0.828610
## V4A48        -2.125e+00  1.239e+00  -1.714 0.086467 .
## V4A49        -9.695e-01  4.179e-01  -2.320 0.020349 *
## V5           1.541e-04  5.538e-05   2.783 0.005385 **
## V6A62        -7.091e-01  3.689e-01  -1.922 0.054561 .
## V6A63        -9.984e-01  5.695e-01  -1.753 0.079581 .
## V6A64        -1.523e+00  6.378e-01  -2.388 0.016959 *
## V6A65        -8.158e-01  3.109e-01  -2.624 0.008692 **
## V7A72        -2.747e-01  5.797e-01  -0.474 0.635653
## V7A73        -9.896e-02  5.532e-01  -0.179 0.858011
## V7A74        -7.037e-01  5.818e-01  -1.209 0.226505
## V7A75        -1.353e-01  5.505e-01  -0.246 0.805883
## V8           4.166e-01  1.128e-01   3.692 0.000222 ***
## V9A92        -1.219e-01  4.772e-01  -0.255 0.798446
## V9A93        -7.147e-01  4.688e-01  -1.525 0.127355
## V9A94        -1.837e-01  5.538e-01  -0.332 0.740136
## V10A102       -3.486e-03  5.052e-01  -0.007 0.994494
## V10A103       -1.457e+00  5.152e-01  -2.828 0.004686 **
## V11           1.035e-02  1.049e-01   0.099 0.921417
## V12A122       6.786e-01  3.209e-01   2.115 0.034449 *
## V12A123       4.051e-01  3.002e-01   1.350 0.177146
## V12A124       1.025e+00  5.380e-01   1.906 0.056681 .
## V13          -1.514e-02  1.144e-02  -1.323 0.185806
## V14A142       -2.349e-01  5.100e-01  -0.461 0.645129
## V14A143       -9.993e-01  2.856e-01  -3.499 0.000467 ***
## V15A152       -6.193e-01  2.985e-01  -2.075 0.038017 *
## V15A153       -9.986e-01  5.829e-01  -1.713 0.086704 .
## V16           2.001e-01  2.219e-01   0.902 0.367146

```

```

## V17A172      3.168e-01  9.030e-01   0.351 0.725716
## V17A173      3.490e-01  8.764e-01   0.398 0.690480
## V17A174      3.083e-01  8.873e-01   0.347 0.728216
## V18          3.457e-01  3.063e-01   1.129 0.259092
## V19A192     -3.343e-01  2.469e-01  -1.354 0.175839
## V20A202     -1.439e+00  8.745e-01  -1.646 0.099772 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 848.32  on 699  degrees of freedom
## Residual deviance: 607.47  on 651  degrees of freedom
## AIC: 705.47
##
## Number of Fisher Scoring iterations: 5

#Create second iteration with all variables with p < 0.05
lreg <- glm(V21~ V1+V2+V3+V4+V5+V6+V8+V10+V12+V14+V15, family=binomial(link =
"logit"), data = train)
summary(lreg)

##
## Call:
## glm(formula = V21 ~ V1 + V2 + V3 + V4 + V5 + V6 + V8 + V10 +
##      V12 + V14 + V15, family = binomial(link = "logit"), data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.1960  -0.7111  -0.3650   0.7109   2.6298
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  5.321e-01  7.313e-01   0.728  0.46686
## V1A12        3.874e-03  2.560e-01   0.015  0.98793
## V1A13       -1.049e+00  4.451e-01  -2.357  0.01844 *
## V1A14       -1.532e+00  2.741e-01  -5.590 2.27e-08 ***
## V2           3.230e-02  1.096e-02   2.947  0.00321 **
## V3A31       -6.505e-01  6.520e-01  -0.998  0.31843
## V3A32       -7.683e-01  5.220e-01  -1.472  0.14105
## V3A33       -1.145e+00  5.819e-01  -1.968  0.04906 *
## V3A34       -1.507e+00  5.471e-01  -2.755  0.00587 **
## V4A41       -1.800e+00  4.467e-01  -4.030 5.59e-05 ***
## V4A410      -2.110e+00  9.760e-01  -2.162  0.03063 *
## V4A42       -8.574e-01  3.089e-01  -2.775  0.00551 **
## V4A43       -9.497e-01  2.890e-01  -3.286  0.00102 **
## V4A44       -4.515e-01  1.256e+00  -0.360  0.71917
## V4A45       -1.150e-01  6.174e-01  -0.186  0.85219
## V4A46       -3.654e-03  4.541e-01  -0.008  0.99358
## V4A48       -1.914e+00  1.189e+00  -1.609  0.10756

```

```

## V4A49      -1.013e+00  4.005e-01  -2.529  0.01145 *
## V5         1.186e-04  5.073e-05   2.339  0.01936 *
## V6A62      -7.112e-01  3.599e-01  -1.976  0.04814 *
## V6A63      -1.069e+00  5.512e-01  -1.939  0.05254 .
## V6A64      -1.446e+00  6.272e-01  -2.306  0.02113 *
## V6A65      -8.754e-01  2.974e-01  -2.943  0.00325 **
## V8         3.194e-01  1.033e-01   3.093  0.00198 **
## V10A102    -1.839e-01  4.901e-01  -0.375  0.70750
## V10A103    -1.454e+00  5.117e-01  -2.841  0.00450 **
## V12A122     5.488e-01  3.054e-01   1.797  0.07233 .
## V12A123     3.835e-01  2.864e-01   1.339  0.18049
## V12A124     8.913e-01  5.049e-01   1.765  0.07754 .
## V14A142    -7.430e-02  4.928e-01  -0.151  0.88017
## V14A143    -8.980e-01  2.756e-01  -3.258  0.00112 **
## V15A152    -7.350e-01  2.712e-01  -2.710  0.00673 **
## V15A153    -1.151e+00  5.427e-01  -2.121  0.03394 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 848.32  on 699  degrees of freedom
## Residual deviance: 627.16  on 667  degrees of freedom
## AIC: 693.16
##
## Number of Fisher Scoring iterations: 5

#Create third iteration with all variables with p < 0.05
lreg <- glm(V21~ V1+V2+V3+V4+V5+V6+V8+V10+V14+V15, family=binomial(link = "logit"), data = train)
summary(lreg)

##
## Call:
## glm(formula = V21 ~ V1 + V2 + V3 + V4 + V5 + V6 + V8 + V10 +
##      V14 + V15, family = binomial(link = "logit"), data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2644  -0.7100  -0.3733   0.7302   2.7433
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  7.691e-01  7.171e-01   1.073  0.283478
## V1A12        -3.122e-02  2.529e-01  -0.123  0.901779
## V1A13        -1.052e+00  4.422e-01  -2.380  0.017334 *
## V1A14        -1.562e+00  2.721e-01  -5.739  9.50e-09 ***
## V2           3.319e-02  1.083e-02   3.065  0.002177 **
## V3A31        -6.128e-01  6.453e-01  -0.950  0.342254
## V3A32        -7.830e-01  5.200e-01  -1.506  0.132089

```

```
## V3A33      -1.130e+00  5.793e-01  -1.951  0.051071 .
## V3A34      -1.506e+00  5.456e-01  -2.760  0.005784 **
## V4A41      -1.775e+00  4.451e-01  -3.988  6.68e-05 ***
## V4A410     -2.168e+00  9.852e-01  -2.201  0.027771 *
## V4A42      -7.973e-01  3.045e-01  -2.619  0.008830 **
## V4A43      -9.597e-01  2.876e-01  -3.337  0.000848 ***
## V4A44      -5.494e-01  1.249e+00  -0.440  0.659893
## V4A45      -1.146e-01  6.051e-01  -0.189  0.849831
## V4A46       1.456e-01  4.453e-01   0.327  0.743697
## V4A48      -1.859e+00  1.185e+00  -1.569  0.116733
## V4A49      -1.025e+00  3.988e-01  -2.570  0.010170 *
## V5         1.352e-04  4.947e-05   2.733  0.006279 **
## V6A62      -6.510e-01  3.559e-01  -1.829  0.067334 .
## V6A63      -1.013e+00  5.524e-01  -1.834  0.066685 .
## V6A64      -1.415e+00  6.152e-01  -2.300  0.021435 *
## V6A65      -8.464e-01  2.946e-01  -2.873  0.004065 **
## V8         3.419e-01  1.022e-01   3.345  0.000822 ***
## V10A102    -9.934e-02  4.787e-01  -0.208  0.835604
## V10A103    -1.531e+00  5.020e-01  -3.051  0.002284 **
## V14A142    -1.973e-01  4.906e-01  -0.402  0.687609
## V14A143    -9.244e-01  2.734e-01  -3.381  0.000723 ***
## V15A152    -7.385e-01  2.675e-01  -2.761  0.005769 **
## V15A153    -6.956e-01  3.802e-01  -1.830  0.067286 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 848.32  on 699  degrees of freedom
## Residual deviance: 631.89  on 670  degrees of freedom
## AIC: 691.89
##
## Number of Fisher Scoring iterations: 5
```

#Address the categorical variables

```
train$V1A13[train$V1 == "A13"] <- 1
train$V1A13[train$V1 != "A13"] <- 0

train$V1A14[train$V1 == "A14"] <- 1
train$V1A14[train$V1 != "A14"] <- 0

train$V3A34[train$V3 == "A34"] <- 1
train$V3A34[train$V3 != "A34"] <- 0

train$V4A41[train$V4 == "A41"] <- 1
train$V4A41[train$V4 != "A41"] <- 0

train$V4A410[train$V4 == "A410"] <- 1
train$V4A410[train$V4 != "A410"] <- 0
```



```

train$V4A42[train$V4 == "A42"] <- 1
train$V4A42[train$V4 != "A42"] <- 0

train$V4A43[train$V4 == "A43"] <- 1
train$V4A43[train$V4 != "A43"] <- 0

train$V4A49[train$V4 == "A49"] <- 1
train$V4A49[train$V4 != "A49"] <- 0

train$V6A64[train$V6 == "A64"] <- 1
train$V6A64[train$V6 != "A64"] <- 0

train$V6A65[train$V6 == "A65"] <- 1
train$V6A65[train$V6 != "A65"] <- 0

train$V10A103[train$V10 == "A103"] <- 1
train$V10A103[train$V10 != "A103"] <- 0

train$V14A143[train$V14 == "A143"] <- 1
train$V14A143[train$V14 != "A143"] <- 0

train$V15A152[train$V15 == "A152"] <- 1
train$V15A152[train$V15 != "A152"] <- 0

lreg <- glm(V21~ V1A13+V1A14+V2+V3A34+V4A41+V4A410+V4A42+V4A43+V4A49+V5+V6A64
+V6A65+V8+V10A103+V14A143+V15A152, family=binomial(link = "logit"), data = tr
ain)
summary(lreg)

##
## Call:
## glm(formula = V21 ~ V1A13 + V1A14 + V2 + V3A34 + V4A41 + V4A410 +
##      V4A42 + V4A43 + V4A49 + V5 + V6A64 + V6A65 + V8 + V10A103 +
##      V14A143 + V15A152, family = binomial(link = "logit"), data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2563  -0.7406  -0.3981   0.8149   2.7237
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -4.529e-01  4.543e-01  -0.997  0.318862
## V1A13        -9.802e-01  4.169e-01  -2.351  0.018702 *
## V1A14        -1.637e+00  2.389e-01  -6.852  7.30e-12 ***
## V2           3.102e-02  1.061e-02   2.925  0.003447 **
## V3A34        -6.940e-01  2.422e-01  -2.865  0.004172 **
## V4A41        -1.715e+00  4.285e-01  -4.002  6.28e-05 ***
## V4A410       -2.311e+00  9.305e-01  -2.484  0.012995 *
## V4A42        -6.579e-01  2.783e-01  -2.364  0.018077 *
## V4A43        -9.249e-01  2.594e-01  -3.566  0.000363 ***

```

```

## V4A49      -8.619e-01  3.516e-01  -2.452 0.014225 *
## V5         1.295e-04  4.787e-05   2.706 0.006813 **
## V6A64      -1.294e+00  6.061e-01  -2.135 0.032766 *
## V6A65      -7.435e-01  2.841e-01  -2.617 0.008870 **
## V8         3.301e-01  9.876e-02   3.343 0.000829 ***
## V10A103    -1.325e+00  4.862e-01  -2.724 0.006445 **
## V14A143    -8.612e-01  2.350e-01  -3.665 0.000247 ***
## V15A152    -4.660e-01  2.161e-01  -2.157 0.031012 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 848.32  on 699  degrees of freedom
## Residual deviance: 649.02  on 683  degrees of freedom
## AIC: 683.02
##
## Number of Fisher Scoring iterations: 5

#Now we can validate
validate$V1A13[validate$V1 == "A13"] <- 1
validate$V1A13[validate$V1 != "A13"] <- 0

validate$V1A14[validate$V1 == "A14"] <- 1
validate$V1A14[validate$V1 != "A14"] <- 0

validate$V3A34[validate$V3 == "A34"] <- 1
validate$V3A34[validate$V3 != "A34"] <- 0

validate$V4A41[validate$V4 == "A41"] <- 1
validate$V4A41[validate$V4 != "A41"] <- 0

validate$V4A410[validate$V4 == "A410"] <- 1
validate$V4A410[validate$V4 != "A410"] <- 0

validate$V4A42[validate$V4 == "A42"] <- 1
validate$V4A42[validate$V4 != "A42"] <- 0

validate$V4A43[validate$V4 == "A43"] <- 1
validate$V4A43[validate$V4 != "A43"] <- 0

validate$V4A49[validate$V4 == "A49"] <- 1
validate$V4A49[validate$V4 != "A49"] <- 0

validate$V6A64[validate$V6 == "A64"] <- 1
validate$V6A64[validate$V6 != "A64"] <- 0

validate$V6A65[validate$V6 == "A65"] <- 1
validate$V6A65[validate$V6 != "A65"] <- 0

```

```
validate$V10A103[validate$V10 == "A103"] <- 1
validate$V10A103[validate$V10 != "A103"] <- 0
```

```
validate$V14A143[validate$V14 == "A143"] <- 1
validate$V14A143[validate$V14 != "A143"] <- 0
```

```
validate$V15A152[validate$V15 == "A152"] <- 1
validate$V15A152[validate$V15 != "A152"] <- 0
```

#Now we can test the model

```
pred <- predict(lreg, validate, type = "response")
pred
```

```
##           2           3           14           15           16           21
## 0.553568685 0.056969561 0.436923371 0.498319732 0.383700733 0.096699687
##           24           28           29           30           34           37
## 0.069632580 0.043206655 0.075218133 0.749356493 0.032501921 0.456897734
##           38           44           47           49           52           53
## 0.337810868 0.159692958 0.167080308 0.075645738 0.174449540 0.056496259
##           65           66           76           79           82           83
## 0.134181850 0.209188508 0.137978581 0.089705674 0.176915514 0.112988732
##           85           87           104          108          113          124
## 0.248895835 0.252509580 0.306271594 0.509923161 0.739923467 0.221618020
##          130          131          133          135          138          144
## 0.341155087 0.597440225 0.149689577 0.376101109 0.286394923 0.288958462
##          150          153          156          161          162          163
## 0.014122647 0.379795360 0.315748809 0.050758173 0.109582947 0.187351775
##          167          168          170          171          172          174
## 0.398297612 0.134001193 0.264992250 0.644484926 0.092546870 0.092990399
##          180          184          187          190          193          197
## 0.394589322 0.029076080 0.362686698 0.433369239 0.505821621 0.012865722
##          201          203          208          210          211          216
## 0.083968239 0.143350305 0.355082964 0.007541972 0.008396107 0.070399330
##          222          232          236          237          240          243
## 0.657427778 0.156246202 0.612393582 0.469502691 0.080018424 0.511795979
##          244          246          247          248          250          251
## 0.030581897 0.114003063 0.037333114 0.394891351 0.093583991 0.075487171
##          252          253          255          257          258          259
## 0.081342578 0.571073991 0.121412272 0.067685268 0.433448364 0.020998082
##          260          265          272          276          279          281
## 0.038263399 0.065956883 0.059624180 0.032834058 0.049215887 0.008481598
##          282          287          295          297          300          301
## 0.101754145 0.483755995 0.257865621 0.016288288 0.081015384 0.056607705
##          310          313          315          316          318          325
## 0.579679650 0.178312060 0.054368299 0.771390606 0.176052635 0.109242185
##          326          327          329          332          336          340
## 0.189141819 0.036893771 0.338830592 0.120233084 0.152990037 0.352009085
##          343          346          352          354          355          357
## 0.282058375 0.044759577 0.153884966 0.563841898 0.226938856 0.008369745
##          358          360          364          366          374          375
```

##	0.209792479	0.643563681	0.053986251	0.045380207	0.684893130	0.881032296
##	378	382	387	388	393	395
##	0.035505682	0.550028403	0.069507961	0.637210156	0.851321767	0.028820779
##	397	402	406	408	409	410
##	0.431587231	0.250701747	0.203228153	0.313733652	0.100867382	0.162493848
##	411	415	420	422	424	427
##	0.404886081	0.430535751	0.375263611	0.092559558	0.130506109	0.076486218
##	428	431	433	434	440	441
##	0.013362698	0.034004847	0.285618627	0.144430355	0.135822283	0.185455928
##	442	443	446	447	452	456
##	0.232253003	0.166746649	0.067407156	0.710091835	0.043193226	0.061847411
##	458	476	484	489	495	500
##	0.226626432	0.336237789	0.022127114	0.187503068	0.408427542	0.080434523
##	502	504	505	512	517	518
##	0.368144390	0.421849485	0.712460058	0.092685110	0.189796051	0.167655688
##	519	520	528	529	534	535
##	0.389905195	0.023761184	0.017015467	0.621643743	0.108448740	0.067998179
##	540	546	548	549	550	553
##	0.218645955	0.715791716	0.175892691	0.482530102	0.056993158	0.213518800
##	557	574	578	579	586	587
##	0.515939685	0.320461536	0.053676564	0.581407083	0.510458820	0.330486378
##	589	590	592	593	595	596
##	0.563774877	0.185674752	0.565129989	0.073387047	0.150248588	0.429954660
##	606	607	611	616	619	624
##	0.701519908	0.065908960	0.502179145	0.769540031	0.516154972	0.490377452
##	628	629	634	635	638	641
##	0.341981125	0.164992996	0.082216384	0.532543633	0.551666700	0.548846145
##	642	646	658	662	664	669
##	0.543693115	0.252545218	0.477427309	0.328666917	0.516462418	0.495490141
##	675	676	681	683	688	690
##	0.247658311	0.150390279	0.062904297	0.181969104	0.816674300	0.130235156
##	696	697	701	704	705	707
##	0.025117060	0.100303830	0.145509326	0.688828692	0.460989530	0.824773754
##	708	709	710	713	720	722
##	0.526326667	0.279079547	0.183743757	0.027063567	0.443640653	0.454085507
##	726	730	732	736	743	747
##	0.024971944	0.017289196	0.359668844	0.723433283	0.062749862	0.451812583
##	748	750	752	755	760	762
##	0.437235997	0.026252798	0.317532252	0.087027601	0.333652796	0.374756398
##	765	768	769	776	781	784
##	0.151034070	0.012216196	0.183167414	0.546136503	0.098461332	0.707278792
##	795	798	799	801	805	806
##	0.241061769	0.042819073	0.118967933	0.202192607	0.404346869	0.768128133
##	811	814	816	821	828	834
##	0.216554766	0.749235287	0.806216968	0.128827433	0.159870746	0.200169351
##	835	838	842	846	856	862
##	0.094645465	0.044088623	0.096164049	0.205083652	0.238547283	0.107300098
##	865	868	870	874	879	881
##	0.183829784	0.079925499	0.451478521	0.094396621	0.534881154	0.084378888
##	883	884	885	894	895	898

```
## 0.335397884 0.046781719 0.455717951 0.209115208 0.023006874 0.011055260
##          901          907          909          911          916          919
## 0.273937631 0.326405896 0.020425417 0.209599295 0.726193588 0.257730009
##          924          925          927          935          936          937
## 0.576381911 0.672178603 0.342741377 0.510757152 0.658349188 0.089554972
##          939          941          942          944          947          950
## 0.935969805 0.064717556 0.128333851 0.041575689 0.628754512 0.078148910
##          951          953          956          957          958          961
## 0.210471125 0.318420118 0.118942558 0.177007424 0.090298475 0.029377798
##          963          968          974          975          976          977
## 0.095880666 0.317273837 0.875562769 0.081837806 0.143401920 0.043705507
##          978          989          990          991          995         1000
## 0.232872176 0.472153586 0.248223312 0.033574263 0.103604189 0.229769893
```

#Threshold is 0.5 here

```
rounded_pred <- as.integer(pred > 0.5)
```

#Create confusion matrix

```
tab <- table(rounded_pred, validate$V21)
```

```
tab
```

```
##
```

```
## rounded_pred    0    1
```

```
##           0 186  57
```

```
##           1  20  37
```

#Calculate accuracy

```
accuracy <- (tab[1,1] + tab[2,2])/sum(tab)
```

```
accuracy
```

```
## [1] 0.7433333
```

#Create ROC curve

```
curve <- roc(validate$V21, rounded_pred)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
accuracy
```

```
## [1] 0.7433333
```

```
curve
```

```
##
```

```
## Call:
```

```
## roc.default(response = validate$V21, predictor = rounded_pred)
```

```
##
```

```
## Data: rounded_pred in 206 controls (validate$V21 0) < 94 cases (validate$V21 1).
```

```
## Area under the curve: 0.6483
```

```
plot(curve, main="Receiver Operating Characteristic (ROC) Curve")
```

Question 10.3.2

#We can iterate through different threshold values to find the best ones

```
loss <- c()

for(x in 1:100) {
  rounded_pred <- as.integer(pred > (x/100))
  tmatrix <- as.matrix(table(rounded_pred, validate$V21))
  if(nrow(tmatrix)>1){
    c1 <- tmatrix[2,1]
  }else{
    c1 <- 0
  }
  if(ncol(tmatrix) > 1){
    c2 <- tmatrix[1,2]
  }else{
    c2 <- 0
  }
  #Perform cost of loss calculations
  loss <- c(loss, c2*5 + c1)
}

plot(c(1:100)/100,loss,xlab = "Threshold Values",ylab = "Loss Values",main =
"Loss vs. Threshold")

loss

## [1] 202 194 182 173 169 160 149 147 154 160 158 154 156 156 157 173 180
174
## [19] 176 176 180 182 184 181 183 190 195 199 207 207 206 219 222 234 243
246
## [37] 250 258 262 265 262 262 272 290 289 296 300 298 302 305 314 332 337
347
## [55] 356 365 374 377 382 381 381 386 396 395 405 409 409 414 412 412 416
431
## [73] 435 434 444 444 448 453 453 453 452 451 456 456 456 455 455 460 465
465
## [91] 465 465 465 470 470 470 470 470 470 470

which.min(loss)

## [1] 8

rounded_pred <- as.integer(pred > (which.min(loss)/100))
tab <- table(rounded_pred, validate$V21)
accuracy <- (tab[1,1] + tab[2,2])/sum(tab)
curve <- roc(validate$V21, rounded_pred)
```

```
## Setting levels: control = 0, case = 1
## Setting direction: controls < cases
accuracy
## [1] 0.5366667
curve
##
## Call:
## roc.default(response = validate$V21, predictor = rounded_pred)
##
## Data: rounded_pred in 206 controls (validate$V21 0) < 94 cases (validate$V
21 1).
## Area under the curve: 0.6568
```