

Question 11.1

For Parts 2 and 3, remember to scale the data first – otherwise, the regression coefficients will be on different scales and the constraint won't have the desired effect.

For Parts 2 and 3, use the `glmnet` function in R.

Notes on R:

- For the elastic net model, what we called λ in the videos, `glmnet` calls “alpha”; you can get a range of results by varying alpha from 1 (lasso) to 0 (ridge regression) [and, of course, other values of alpha in between].
- In a function call like `glmnet(x,y,family="mgaussian",alpha=1)` the predictors `x` need to be in R's matrix format, rather than data frame format. You can convert a data frame to a matrix using `as.matrix` – for example, `x <- as.matrix(data[,1:n-1])`
- Rather than specifying a value of `T`, `glmnet` returns models for a variety of values of `T`.


Using the crime data set `uscrime.txt` from Questions 8.2, 9.1, and 10.1, build a regression model using:

1. Stepwise regression

The code for this approach can be found in [appendix 11.1.1](#). In order to approach the stepwise regression, I first began by scaling the data. Once that was done, I wanted to create the stepwise regression model, so I used the `stepAIC` function with `direction` set to both, so it would do a backward and forward search. In doing so, the predictors that the model identified were: `M`, `Ed`, `Po1`, `M.F`, `U1`, `U2`, `Ineq`, `Prob`. Then, I wanted to identify the best number of predictors. For this, I set up a 10-fold cross-validation approach, trained the model on the scaled data, and used the “`leapSeq`” method, which would fit the linear regression with stepwise selection. I also set `nvmax` values from 1:15 so that it would identify the best models with 1 predictor, 2 predictors, ..., 15 predictors. Upon finishing this process, I was able to identify that the best model was one with six predictors as it had the lowest RMSE (232.8635). The predictors present in this model were `M`, `Ed`, `Po1`, `U2`, `Ineq`, and `Prob`. I then wanted to cross-validate and calculate the R-squared values for the model with the originally identified predictors and the model with only six predictors. In the end, the R-squared value for the model with predictors: `M`, `Ed`, `Po1`, `M.F`, `U1`, `U2`, `Ineq`, and `Prob` was 0.667621. The R-squared value for the model with predictors: `M`, `Ed`, `Po1`, `U2`, `Ineq`, and `Prob` was 0.6661638. Ultimately, both values were pretty similar here, and it seems that `M.F` and `U1` were not too significant overall.

2. Lasso

The code for this approach can be found in [appendix 11.1.2](#). In order to approach lasso regression, I also began by first scaling the data. Once that was done, I used the `cv.glmnet` function to perform the actual lasso regression and plotted the output:



The above graph displays the MSE against the log of lambda. The leftmost vertical line shows that the log of the optimal value of lambda is approximately two. In particular, the value of lambda is 9.237784. Following this, I wanted to find the non-zero coefficient variables of the model associated with the minimum (optimal) value of lambda. They are: So, M, Ed, Po1, M.F, NW, U1, U2, Wealth, Ineq, and Prob. I then fitted a new model with the aforementioned factors. Upon doing this, it showed me there were only six significant variables: M, Ed, Po1, U2, Ineq, and Prob. I then wanted to cross-validate and calculate the R-squared value of the original model and the model with only the six significant variables. The R-squared value for the original model, with predictors: So, M, Ed, Po1, M.F, NW, U1, U2, Wealth, Ineq, and Prob, was 0.6051577. The R-squared value for the model with only the significant predictors: M, Ed, Po1, U2, Ineq, and Prob was 0.6661638. Here, the R-squared value increased as we removed the predictors So, M.F, NW, U1, and Wealth, indicating that they are likely not essential to the model.

3. Elastic net

The code for this approach can be found in [appendix 11.1.3](#). In order to approach the elastic net regression, I first began by scaling the data. Then, I wanted to vary the alpha values from 0 to 1, incrementing by 0.1 each time, to find the best value. In order to identify which was best, I looked at which alpha value resulted in the highest R-squared value among those iterated through. As a result of this, I found the best alpha value was 0.9. Once this was done, I built an elastic net model using the best value of alpha. The important variables identified were: So, M, Ed, Po1, M.F, Pop, NW, U1, U2, Wealth, Ineq, and Prob. I then used these variables to fit a new model. From this, there were six significant variables identified: M, Ed, Po1, U2, Ineq, and Prob.

Then, I used cross-validation to calculate the R-squared value of the original model and the model with only six significant variables. The R-squared value for the original model, with predictors: So, M, Ed, Po1, M.F, Pop, NW, U1, U2, Wealth, Ineq, and Prob, was 0.5903894. The R-squared value for the model with only the six significant predictors: M, Ed, Po1, U2, Ineq, and Prob was 0.6661638. Here, the R-squared values once again increased as we removed inessential variables such as So, M.F, Pop, NW, U1, and Wealth.

Summary

Model	Relevant Variables	CV R-Squared
Stepwise Regression, All Variables	M, Ed, Po1, M.F, U1, U2, Ineq, and Prob	0.667621
Stepwise Regression, Significant Variables	M, Ed, Po1, U2, Ineq, and Prob	0.6661638
LASSO Regression, All Variables	So, M, Ed, Po1, M.F, NW, U1, U2, Wealth, Ineq, and Prob	0.6051577
LASSO Regression, Significant Variables	M, Ed, Po1, U2, Ineq, and Prob	0.6661638
Elastic Net, All Variables	So, M, Ed, Po1, M.F, Pop, NW, U1, U2, Wealth, Ineq, and Prob	0.5903894
Elastic Net, Significant Variables	M, Ed, Po1, U2, Ineq, and Prob	0.6661638

There are a few interesting observations here. First of all, it appears that, when using all variables, elastic net regression performed the worst out of all three while stepwise regression performed the best. It is also interesting to note that elastic net, when using all variables, had the most at 12 while LASSO had 11, and stepwise had 8. Therefore, decreasing the number of variables selected does seem to improve the overall quality of the model. Furthermore, for each model, when I looked for the significant variables identified, the same six were identified each time: M, Ed, Po1, U2, Ineq, and Prob. In addition to this, more often than not, the six variables also produced the highest cross-validation R-squared value when fitted to a model (0.6661638) when compared to the R-squared value produced by fitting all identified variables. It is also likely that these variable-selection models were particularly helpful in this situation because there was a small number of data points. As per lecture, in these situations, we do not want to have too many factors as it can cause overfitting. Overall, it seems that variable-selection models are extremely helpful and can produce really good results, especially compared to some of the other models we have used in class.

Appendix

Question 11.1.1

#This is the stepwise regression approach

#Load Packages

```
library(tidyverse); library(caret); library(leaps); library(MASS)
```

```
## — Attaching packages ————— tidyverse 1.3.2 —
```

```
## ✓ ggplot2 3.4.0      ✓ purrr  1.0.1
## ✓ tibble  3.1.8      ✓ dplyr  1.0.10
## ✓ tidyr   1.2.1      ✓ stringr 1.5.0
## ✓ readr   2.1.3      ✓ forcats 0.5.2
```

```
## — Conflicts ————— tidyverse_conflicts() —
```

```
## ✗ dplyr::filter() masks stats::filter()
```

```
## ✗ dplyr::lag()     masks stats::lag()
```

```
## Loading required package: lattice
```

```
##
```

```
##
```

```
## Attaching package: 'caret'
```

```
##
```

```
##
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## lift
```

```
##
```

```
##
```

```
##
```

```
## Attaching package: 'MASS'
```

```
##
```

```
##
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
## select
```

#Load Data

```
data <- read.table("C:\\Users\\User\\OneDrive\\Desktop\\Data 11.1\\uscrime.txt", stringsAsFactors = F, header = T)
```

#Scale Data

```
scaledD <- as.data.frame(scale(data[,c(1,3,4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15)]))
```

```
scaledD <- cbind(data[,2], scaledD, data[,16])
```

```
colnames(scaledD)[1] <- "So"
```

```
colnames(scaledD)[16] <- "Crime"
```

#Stepwise regression

```

data_model <- lm(Crime~., data = scaledD)
step_model <- stepAIC(data_model, direction = "both", trace = F)
summary(step_model)

##
## Call:
## lm(formula = Crime ~ M + Ed + Po1 + M.F + U1 + U2 + Ineq + Prob,
##     data = scaledD)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -444.70 -111.07   3.03  122.15  483.30
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   905.09      28.52  31.731 < 2e-16 ***
## M              117.28      42.10   2.786  0.00828 **
## Ed             201.50      59.02   3.414  0.00153 **
## Po1            305.07      46.14   6.613 8.26e-08 ***
## M.F             65.83      40.08   1.642  0.10874
## U1            -109.73      60.20  -1.823  0.07622 .
## U2             158.22      61.22   2.585  0.01371 *
## Ineq           244.70      55.69   4.394 8.63e-05 ***
## Prob          -86.31      33.89  -2.547  0.01505 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 195.5 on 38 degrees of freedom
## Multiple R-squared:  0.7888, Adjusted R-squared:  0.7444
## F-statistic: 17.74 on 8 and 38 DF, p-value: 1.159e-10

#Identifying best number of predictors
set.seed(1)

#10-fold CV
train.control <- trainControl(method = "cv", number = 10)

#Training model
step.model <- train(Crime~., data = scaledD, method = "leapSeq", tuneGrid = d
ata.frame(nvmax = 1:15), trControl = train.control)

step.model$results

```

	nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	1	271.2008	0.5843691	226.2581	110.47396	0.3371720	96.68780
## 2	2	254.3833	0.6064511	202.4499	84.83446	0.2805282	64.11189
## 3	3	238.7279	0.7033882	186.3716	88.72517	0.1011618	59.18009
## 4	4	257.3026	0.7051499	200.2685	100.10472	0.1940803	76.05679
## 5	5	252.0782	0.6269684	201.9099	91.52018	0.2311683	70.23161
## 6	6	232.8635	0.7854913	184.9587	102.32196	0.1476856	81.55553

```
## 7      7 240.1809 0.7294113 190.5553 95.43075 0.1820599 74.00488
## 8      8 258.5924 0.6986070 202.8368 111.99607 0.2032220 90.26335
## 9      9 260.8053 0.6076826 210.0503 108.82448 0.1764839 89.03990
## 10     10 267.2403 0.6709714 216.7676 96.06399 0.2069373 76.76546
## 11     11 240.1841 0.6456568 201.6880 85.04252 0.2280547 71.83940
## 12     12 273.1258 0.6852765 222.0306 116.50459 0.2130656 93.46969
## 13     13 246.1732 0.5977355 199.2895 85.79255 0.2451306 69.47176
## 14     14 245.9982 0.5961171 195.7384 90.27009 0.2680203 68.68076
## 15     15 247.1168 0.5898489 196.0868 87.86030 0.2573792 67.76324
```

```
step.model$bestTune
```

```
## nvmax
## 6      6
```

```
summary(step.model$finalModel)
```

```
## Subset selection object
## 15 Variables (and intercept)
##      Forced in Forced out
## So      FALSE      FALSE
## M        FALSE      FALSE
## Ed        FALSE      FALSE
## Po1       FALSE      FALSE
## Po2       FALSE      FALSE
## LF        FALSE      FALSE
## M.F       FALSE      FALSE
## Pop       FALSE      FALSE
## NW        FALSE      FALSE
## U1        FALSE      FALSE
## U2        FALSE      FALSE
## Wealth    FALSE      FALSE
## Ineq      FALSE      FALSE
## Prob      FALSE      FALSE
## Time      FALSE      FALSE
```

```
## 1 subsets of each size up to 6
```

```
## Selection Algorithm: 'sequential replacement'
```

```
##      So M Ed Po1 Po2 LF M.F Pop NW U1 U2 Wealth Ineq Prob Time
## 1 ( 1 ) " " " " " " "*" " " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " "*" " " " " " " " " " " " " " " " "
## 3 ( 1 ) " " " " "*" "*" " " " " " " " " " " " " " " " "
## 4 ( 1 ) " " "*" "*" "*" " " " " " " " " " " " " " " " "
## 5 ( 1 ) " " "*" "*" "*" " " " " " " " " " " " " " " " "
## 6 ( 1 ) " " "*" "*" "*" " " " " " " " " " " " " "*" " " "
## 6 ( 1 ) " " "*" "*" "*" " " " " " " " " " " " " "*" " " "
## 6 ( 1 ) " " "*" "*" "*" " " " " " " " " " " " " "*" " " "
```

```
coef(step.model$finalModel, 6)
```

```
## (Intercept)          M          Ed          Po1          U2          Ineq
## 905.08511    131.98475    219.79230    341.84009    75.47364    269.90968
## Prob
## -86.44225
```

```

#Check CV with LOOC
#Using all variables
TotSS <- sum((data$Crime - mean(data$Crime))^2)
TotSSE <- 0

for (i in 1:nrow(scaledD)){
  fit_step_i = lm(Crime ~ M+Ed+Po1+M.F+U1+U2+Ineq+Prob, data=scaledD[-i,])
  pred_i <- predict(fit_step_i, newdata=scaledD[i,])
  TotSSE <- TotSSE + ((pred_i - data[i, 16])^2)
}

R2_all <- 1 - TotSSE/TotSS
R2_all

##          1
## 0.667621

#Using only significant variables
TotSS <- sum((data$Crime - mean(data$Crime))^2)
TotSSE <- 0

for (i in 1:nrow(scaledD)){
  fit_step_i = lm(Crime ~ M+Ed+Po1+U2+Ineq+Prob, data=scaledD[-i,])
  pred_i <- predict(fit_step_i, newdata=scaledD[i,])
  TotSSE <- TotSSE + ((pred_i - data[i, 16])^2)
}

R2_sig <- 1 - TotSSE/TotSS
R2_sig

##          1
## 0.6661638

```

Question 11.1.2

```

#This is the lasso regression approach

#Load Packages
library(tidyverse); library(caret); library(leaps); library(MASS); library(glmnet)

## Loading required package: Matrix

##
## Attaching package: 'Matrix'

## The following objects are masked from 'package:tidyr':
##
##      expand, pack, unpack

## Loaded glmnet 4.1-6

```

```

#Load Data
data <- read.table("C:\\Users\\User\\OneDrive\\Desktop\\Data 11.1\\uscrime.txt", stringsAsFactors = F, header = T)

#Scale Data
scaledD <- as.data.frame(scale(data[,c(1,3,4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15)]))
scaledD <- cbind(data[,2], scaledD, data[,16])
colnames(scaledD)[1] <- "So"
colnames(scaledD)[16] <- "Crime"

set.seed(1)

#Lasso regression

lasso <- cv.glmnet(x=as.matrix(scaledD[, -16]), y=as.matrix(scaledD$Crime), alpha=1, nfolds = 10, type.measure="mse", family = "gaussian")

plot(lasso)

lasso$lambda.min

## [1] 9.237784

#Use value of lambda that gives min. CVM
coef(lasso, s=lasso$lambda.min)

## 16 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) 889.944466
## So          44.475632
## M           89.288569
## Ed          138.032503
## Po1         305.238076
## Po2         .
## LF          .
## M.F         55.156773
## Pop         .
## NW          6.192810
## U1          -36.329654
## U2          71.884874
## Wealth      4.808914
## Ineq        191.893564
## Prob       -83.561850
## Time       .

#Fit new model with new variables
fit_lasso <- lm(Crime ~ So+M+Ed+Po1+M.F+NW+U1+U2+Wealth+Ineq+Prob, data=scaledD)

summary(fit_lasso)

```



```
##
## Call:
## lm(formula = Crime ~ So + M + Ed + Po1 + M.F + NW + U1 + U2 +
##      Wealth + Ineq + Prob, data = scaledD)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -408.38  -96.14   -1.39   114.80   454.53
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   893.73      51.33   17.411 < 2e-16 ***
## So             33.35     123.69    0.270  0.78905
## M             114.97     48.92    2.350  0.02454 *
## Ed            195.31     62.52    3.124  0.00357 **
## Po1           275.69     59.99    4.596 5.41e-05 ***
## M.F           64.50     42.82    1.506  0.14101
## NW            15.93     57.16    0.279  0.78209
## U1           -94.61     64.90   -1.458  0.15380
## U2            140.81     66.32    2.123  0.04089 *
## Wealth         73.59     93.96    0.783  0.43878
## Ineq          267.01     80.66    3.310  0.00217 **
## Prob          -87.64     40.25   -2.177  0.03627 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 201.3 on 35 degrees of freedom
## Multiple R-squared:  0.794, Adjusted R-squared:  0.7292
## F-statistic: 12.26 on 11 and 35 DF, p-value: 5.334e-09

#Check CV with LOOC
#Using all variables
TotSS <- sum((data$Crime - mean(data$Crime))^2)
TotSSE <- 0

for (i in 1:nrow(scaledD)){
  fit_lasso_i = lm(Crime ~ So+M+Ed+Po1+M.F+NW+U1+U2+Wealth+Ineq+Prob, data=scaledD[-i,])
  pred_i <- predict(fit_lasso_i, newdata=scaledD[i,])
  TotSSE <- TotSSE + ((pred_i - data[i, 16])^2)
}

R2_all <- 1 - TotSSE/TotSS
R2_all

##      1
## 0.6051577

#Using only significant variables
TotSS <- sum((data$Crime - mean(data$Crime))^2)
```

```
TotSSE <- 0

for (i in 1:nrow(scaledD)){
  fit_lasso_i = lm(Crime ~ M+Ed+Po1+U2+Ineq+Prob, data=scaledD[-i,])
  pred_i <- predict(fit_lasso_i, newdata=scaledD[i,])
  TotSSE <- TotSSE + ((pred_i - data[i, 16])^2)
}

R2_sig <- 1 - TotSSE/TotSS
R2_sig

##          1
## 0.6661638
```

Question 11.1.3

#This is the elastic net regression approach

#Load Packages

```
library(tidyverse); library(caret); library(leaps); library(MASS); library(glmnet)
```

#Load Data

```
data <- read.table("C:\\Users\\User\\OneDrive\\Desktop\\Data 11.1\\uscrime.txt", stringsAsFactors = F, header = T)
```

#Scale Data

```
scaledD <- as.data.frame(scale(data[,c(1,3,4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15)]))
```

```
scaledD <- cbind(data[,2], scaledD, data[,16])
```

```
colnames(scaledD)[1] <- "So"
```

```
colnames(scaledD)[16] <- "Crime"
```

```
set.seed(1)
```

#Vary alpha values, calculate R-squared

```
R2_el <- c()
```

```
for (i in 0:10){
```

```
  fit_elastic <- cv.glmnet(x=as.matrix(scaledD[, -16]), y=as.matrix(scaledD$Crime), alpha = i/10, nfolds = 10, type.measure="mse", family="gaussian")
```

```
  R2_el <- cbind(R2_el, fit_elastic$glmnet.fit$dev.ratio[which(fit_elastic$glmnet.fit$lambda==fit_elastic$lambda.min)])
```

```
}
```

```
R2_el
```

```
##          [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.7450772 0.7490443 0.7446379 0.7167649 0.7853621 0.7936385 0.7597881
##          [,8]      [,9]     [,10]     [,11]
## [1,] 0.7940936 0.7705495 0.7941416 0.7682889
```

```

#The best alpha value is:
best_alpha <- (which.max(R2_el)-1)/10
best_alpha

## [1] 0.9

#Build model using best alpha
elastic <- cv.glmnet(x=as.matrix(scaledD[, -16]), y=as.matrix(scaledD$Crime),
alpha = best_alpha, nfolds = 10, type.measure = "mse", family="gaussian")

coef(elastic, s=elastic$lambda.min)

## 16 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) 894.24787
## So          31.83438
## M           104.63624
## Ed          175.94404
## Po1         295.00439
## Po2         .
## LF          .
## M.F         52.69543
## Pop        -19.58677
## NW          15.25481
## U1         -72.62608
## U2          117.92638
## Wealth      55.75815
## Ineq        251.09728
## Prob       -89.64831
## Time        .

#Use important variables to create new model

elastic_imp <- lm(Crime~So+M+Ed+Po1+M.F+Pop+NW+U1+U2+Wealth+Ineq+Prob, data=scaledD)

summary(elastic_imp)

##
## Call:
## lm(formula = Crime ~ So + M + Ed + Po1 + M.F + Pop + NW + U1 +
##      U2 + Wealth + Ineq + Prob, data = scaledD)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -434.18 -107.01   18.55  115.88  470.32
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   897.29      51.91  17.286 < 2e-16 ***
## So            22.89      125.35   0.183  0.85621

```

```

## M            112.71      49.35    2.284  0.02876 *
## Ed           195.70      62.94    3.109  0.00378 **
## Po1          293.18      64.99    4.511  7.32e-05 ***
## M.F          48.92      48.12    1.017  0.31656
## Pop          -33.25      45.63   -0.729  0.47113
## NW           19.16      57.71    0.332  0.74195
## U1           -89.76      65.68   -1.367  0.18069
## U2           140.78      66.77    2.108  0.04245 *
## Wealth       83.30      95.53    0.872  0.38932
## Ineq         285.77      85.19    3.355  0.00196 **
## Prob        -92.75      41.12   -2.255  0.03065 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 202.6 on 34 degrees of freedom
## Multiple R-squared:  0.7971, Adjusted R-squared:  0.7255
## F-statistic: 11.13 on 12 and 34 DF,  p-value: 1.52e-08

#Check CV with LOOC
#Using all variables
TotSS <- sum((data$Crime - mean(data$Crime))^2)
TotSSE <- 0

for (i in 1:nrow(scaledD)){
  fit_elastic_i = lm(Crime~So+M+Ed+Po1+M.F+Pop+NW+U1+U2+Wealth+Ineq+Prob, data=scaledD[-i,])
  pred_i <- predict(fit_elastic_i, newdata=scaledD[i,])
  TotSSE <- TotSSE + ((pred_i - data[i, 16])^2)
}

R2_all <- 1 - TotSSE/TotSS
R2_all

##          1
## 0.5903894

#Using only significant variables
TotSS <- sum((data$Crime - mean(data$Crime))^2)
TotSSE <- 0

for (i in 1:nrow(scaledD)){
  fit_elastic_i = lm(Crime~M+Ed+Po1+U2+Ineq+Prob, data=scaledD[-i,])
  pred_i <- predict(fit_elastic_i, newdata=scaledD[i,])
  TotSSE <- TotSSE + ((pred_i - data[i, 16])^2)
}

R2_sig <- 1 - TotSSE/TotSS
R2_sig

##          1
## 0.6661638

```

